

# Project 2: BitTorrent Client Implementation

Full Project Due: Oct. 16th at 11:00 PM  
Milestone: Sept. 25th at 11pm

## Overview

In this lab, you and your lab partner will develop a basic BitTorrent client that can exchange a file between multiple peers. The programming in this assignment will be in C/C++ requiring standard sockets, and thus you **will not** need to have root access. You should be able to develop your code on any CS lab machine or your personal machine, but all code will be tested on the CS lab.

You should not use external libraries (such as Boost). If you have questions about if a library you are wanting to use is allowed, please ask!

## Readings

The following readings will help you understand the concepts of this project or be helpful in implementing your code. Please remember that any readings assigned here will be material that you should have a good grasp on for assignments or tests!

- Kurose and Ross: Chapter 2
- Donahoo and Calvert (Sockets book): Chapter 2 (TCP sockets), Chapter 5 (Sending Data), Chapter 6 (Advanced Socket Programming)

## Deliverable

Your submission should minimally include the following programs and files (or similarly structured files):

- README
- Makefile
- `bt_lib.c|h`
- `bt_setup.c|h`
- `bt_client`
- `mylog.log` A log file containing connection details and file transfer information (at a minimum) from one of your own tests.

Your README file should contain a short header containing your name, username, and the assignment title. The README should additionally contain a short description of your code, tasks accomplished, and how to compile, execute, and interpret the output of your programs. **Your README should also contain a list of all submitted files and code and the functionality implemented in different code files.** This is a large assignment, and this will greatly aid in grading.

Even though you are working in groups on this assignment, each student must submit their own `write-up`. You may not see your partners written answers. The `write-up` PDF file should contain a short header containing your name, username, and the assignment title. It should also contain answers to all short questions found in this lab as well as the general information about the contents of each file that you submit. That is, a person not familiar with this assignment should have a clear understanding of the purpose of each submitted file after reading your `write-up`. Finally, the `write-up` should include a Credits: section, which lists all the people (e.g., students in the class or not, friends, colleagues, etc.) and resources (e.g., webpages, manuals, books, code, etc.) that you consulted for this assignment.

Your program should always exit gracefully. In addition, the code should be clean and well-commented. Most importantly, the submitted code must compile and run on the CS Silo/Burrow machines.

## **Milestones**

The entire assignment is due on October 16th, but your group will also submit a Milestone within one week of the project start to show you have read the assignment and are on track to finish the project. More details about what is required for the Milestone can be found later in this document.

## **Submission Instructions**

Submit a tarball with your files and folders named `bt_client` via `oncourse`.

## **Cheating and Plagiarism**

Although it is okay to discuss assignments (verbally; no written exchanges) with other students, all solutions handed in must be your own. Your code must be a group effort between the two group members only (you are expected to write the code together). The `write-up` must be written individually; you may not read your project partners (or any other students) `write-up`.

Copying answers from friends or the Internet is prohibited. We will use different software packages to check for plagiarism. Furthermore, you must cite all sources used for problem solutions (i.e. websites, books, discussions with classmates, etc.) in the `write-up`. For more information please see Indiana University Code of Student Rights, Responsibilities, and Conduct: <http://www.indiana.edu/code/>.

# 1 BitTorrent Protocol

BitTorrent is a peer-to-peer file sharing protocol developed by Bram Cohen. It was developed in response to centralized file sharing services, like Napster, where centralized infrastructure was necessary to enable downloaders to connect with uploaders. Additionally, centralized services did not enable for multiple file uploaders to send data to multiple file downloaders.

BitTorrent is designed to address both of these shortcomings by enabling file downloaders and file hosters to interact directly in a peer-to-peer fashion; that is, BitTorrent participants interact using a peer-to-peer protocol where each participant acts as both a server and client. Perhaps BitTorrent's most important feature is that it enables a large number of downloaders to collaborate in downloading a file from a much smaller number of file uploaders. BitTorrent also has built in mechanisms to ensure fairness which encourages participants to share previously downloaded data as uploaders, and if they don't, a downloader is punished by limiting download speeds.

In this lab, you will implement a core subset of the BitTorrent protocol by implementing a BitTorrent client. Below, I outline the relevant portion of the BitTorrent protocol, for more detailed descriptions, you should read Bram Cohen original specification here: [http://www.bittorrent.org/beps/bep\\_0003.html](http://www.bittorrent.org/beps/bep_0003.html). Unless otherwise stated in this document or clarified on the Project wiki, you are expected to implement your client to meet the specification outlined by Cohen.

## 1.1 BitTorrent Basics

There are three distinct parts of the BitTorrent protocol:

- **Metainfo File:** The *metainfo file* (or a *torrent file*) is a description of the file[s] to be downloaded as well as information to contact the tracker server. A metainfo file is generated by the *file originator* and uploaded to a web server where file downloaders can retrieve the metainfo file and use it to learn about the tracker server.
- **Tracker Server:** The tracker server monitors the number of clients downloading/uploading files. When a file downloader contacts the tracker server, a list of other clients ip and ports are returned. Tracker servers have become rather advanced, and some even actively participate in the BitTorrent client procedure to better monitor and ensure fairness. (*You are not required to implement the tracker server for BitTorrent in this project.*)
- **Clients:** A BitTorrent client is the application run by downloaders/uploaders that exchanges *pieces* of a file. The division of the file into pieces is described in the metainfo file, and the protocol for exchanging file pieces is called the *client protocol* or *peer protocol*. Clients learn about other BitTorrent clients via the `tracker server`, as well as accepting new request from peer clients directly in a peer-to-peer fashion.

There are three distinctions of BitTorrent clients: A client that is only uploading a file is described as a *seeder*; a client that is only downloading a file is described as a *leecher*; and a client that does both, either downloading and/or uploading, are described as *peers*. All seeders, leechers, and peers actively downloading/uploading are described as the *swarm*.

The BitTorrent workflow is as follows. A file originator who wants to start sharing a file generates a metainfo file which contains details about the tracker server, the size and partitioning of the file, and other

relevant details. The metainfo file is uploaded to a generic web server where others who wish to download the file can retrieve the metainfo.

With the metainfo in hand, a BitTorrent client contacts the tracker server described in the metainfo file and retrieves a list of other clients participating in the swarm. The client then contacts the other clients, establishes a connection, and starts the client protocol by requesting and downloading pieces of the file *in random order* until the entire file is downloaded. As the client collects pieces, it starts making these available to others in the swarm by announcing that it has the piece. In this way, together, all clients can aid each other in downloading the file faster and from dispersed sources.

In this lab, your client ***must*** be able to parse and understand metainfo files and participate in the client protocol. However, ***you are not required*** to implement features of the tracker nor have your client contact a tracker to retrieve a list of peer clients. Instead, your BitTorrent client will take a pre-selected list of peers as a command line argument.

## 1.2 Metainfo Files and bencoding

Metadata in BitTorrent is encoded in a special format call *bencoding*. Since we are dealing with an application layer protocol, it is important to implement the setup required for the protocol to work – e.g. reading `.torrent` files – so that you both understand how application protocols may start and the effort that is required to make application layer protocols function. As a result, we will require you to parse `.torrent` files for your first milestone, which we discuss in more detail at the end of this document.

Most relevant to this discussion is that bencoding allows for standard data types, like integers, strings, dictionaries and lists, and a metainfo file is just a bencoded dictionary with the following keys and values:

- `announce` : The URL of the torrent tracker
- `info` : A dictionary with the following keys
  - `name` : The suggested name for saving the file
  - `piece length` : The size, in bytes, of a piece for this file. This is always a power of 2.
  - `length` : The length of the file to be downloaded
  - `pieces` : A string of length in multiples of 20 bytes. Each 20 bytes subpart represents the SHA1 hash of the piece of the file at that index.

There are other formats for downloading multiple-files, but ***you are not required to handle multi-file torrents*** in this assignment. If you are interested, I refer you to the BitTorrent specification. Most importantly, you should think carefully about the `info` portion of the metainfo file. Particularly, the `pieces` value which not only describes the data expectation of the file (i.e., its hash), but also indicates the index of the pieces. The client protocol is all about exchanging pieces; specifically, the client protocol is concerned with the requesting, announcing, and downloading of pieces.

The following website may be helpful in understanding how bencoding works and structures its information: <https://wiki.theory.org/BitTorrentSpecification#Bencoding>.

## 1.3 Tracker Protocol

Since you are not required to implement or communicate with a tracker, I refer you to the BitTorrent specification for details on communicating with the tracker servers.

## 1.4 Client Protocol

The client protocol describes the process by which peers contact each other, learn which pieces others have, and request, download and upload pieces per these requests. Peers establish a TCP connection between each other, and these connections are symmetric, used for both uploading and downloading of pieces between peer clients.

**Connection State** For each connection to a peer, a client maintains two bits of state. A connection can either be *choked* or *unchoked*; that is, a client is willing to upload to an unchoked connection but not to a choked connection. The process of choking a connection is what encourages fairness in BitTorrent as well as limits the amount of uploading a client is willing to do for a given downloader. The other bit of state is based on *interest*, which indicates the client is interested in some piece that the peer has. ***A transfer can move forward when the connection isn't choked and one side is interested; initially, connections are set to choked and not interested.***

**Connection Establishment and Handshakes** Before any data transfer, two peers must first establish a BitTorrent connection via a *handshake*. The handshake protocol is as follows and is symmetric, that is, both sides send the same information.

1. 20 bytes: The first byte is the ASCII character 19 (decimal), i.e., the number 19 encoded in an unsigned byte, and the remaining bytes is the string value "BitTorrent Protocol".
2. 8 bytes: Reserved bytes all set to zero that may be used in extensions of the BitTorrent protocol.
3. 20 byte: SHA1 hash of the bencoded `info` value in the metafile, also known as the `info_hash`. Both sides should send the same value, if not, then they are not interested in participating in the same swarm.
4. 20 byte: The peer ID used to identify this client. ***You can use shorter peer id's but you should still place them within a 20 byte value.*** If the peer ID doesn't match the expected ID, the connection is severed.

**Messages** Once the handshake is complete, further communication is conducted via an exchanges of *messages*. BitTorrent messages can be described in the C structures presented in Figure 1, which I have provided for you in the skeleton code. Below, using Figure 1 as a guide, I describe the components and usage of the key messages.

Starting with the `bt_msg_t` type, a message starts with a `length` which indicates the length of the message, and following, is a one byte field, the `bt_type`, which describes the type of the message. There are different types of messages, described in the `payload union`. Below, we describe each of the types and there type identifiers:

- 0 - `choke`: A message indicating to peer that they are now choked. This message has 1 as its length value.
- 1 - `unchoke`: A message indicating to peer that they are now unchoked. For this project, once you unchoke a peer, they will not need to re-choke them. This message has 1 as its length value.

```

typedef struct {
    char * bitfield; //bitfield where each bit represents a piece that
                        //the peer has or doesn't have
    size_t size; //size of the bitfield
} bt_bitfield_t;

typedef struct{
    int index; //which piece index
    int begin; //offset within piece
    int length; //amount wanted, within a power of two
} bt_request_t;

typedef struct{
    int index; //which piece index
    int begin; //offset within piece
    char piece[0]; //pointer to start of the data for a piece
} bt_piece_t;

typedef struct bt_msg{
    int length; //prefix length, size of remaining message
                //0 length message is a keep-alive message
    unsigned char bt_type; //type of bt_message

    //payload can be any of these
    union {
        bt_bitfield_t bitfield; //send a bitfield
        int have; //index of piece just completed
        bt_piece_t piece; //a piece message
        bt_request_t request; //request message
        bt_request_t cancel; //cancel message, same type as request
        char data[0]; //pointer to start of payload, just for convenience
    } payload;
} bt_msg_t;

```

Figure 1: The basic data structures for BitTorrent's message protocol.

- 2 - interested: A message indicating to peer that they are interested in some piece that at the peer. This message has 1 as its length value.
- 3 - not interested: A message indicating to peer that they are not interested in *any* pieces at the peer. This message has 1 as its length value, the length of `bt_type`.
- 4 - have: A message sent once to all peers once the client has downloaded a completed piece and verified it using the hash in the metainfo file. The value is just an integer indicating the piece index.
- 5 - bitfield: A message *only* sent after the handshake is complete to tell peers which pieces the client has available. The first byte in the bitfield represents pieces indexed 0-7, the second byte represents pieces index 8-15, and so forth. Extra bits are set to 0.
- 6 - request: A message sent to a peer indicating that they would like to download a given piece. A request message contains an `index`, `begin`, and `length` field (see structure above). The `index` indicates which piece is being requested. The `begin` field is the byte offset within the piece, and the `length` is the amount of data requested, in bytes. Generally, the `length` is a power of 2 and must

not exceed  $2^{17}$ . Most implementations use  $2^{15}$  as the length. The subset of the piece request is referred to as the *piece block*.

- 7 - `piece`: A message that indicates a data transfer of a piece (or a block of a piece). It has two fields: `index`, indicating which piece this is associated with; and, `begin`, indicating the byte offset within the piece. The final value, `piece` is a pointer pointing to the start of the raw file data.
- 8 - `cancel`: A cancel message is used to cancel a request for a piece (or block of a piece). A cancel message has the same data layout as a request, and it is generally used in the “End Game” where the file download is nearly complete. At the end of the download, the final blocks and pieces have a tendency to take a long time, so a client sends a request for the block or piece to every peer, and once it is downloaded, it must cancel the request using a cancel message.

Additionally, clients check the live-ness of their open connections by sending keep-alive messages. These messages have 0 set as the message length and are ignored by receivers with respect to taking actions. Keep-alive messages are sent once every two minutes, but if you are already exchanging pieces with a client, you can consider the client alive. ***You must implement a strategy for handling client failures.***

For this assignment, you do not need to implement 0, 3, 8, or keep-alive messages. If you find them useful, however, you may include them.

## 2 Provided Code

We have provided a good portion of skeleton code for you to work from, and remember...you can change ANYTHING in the skeleton code, as the skeleton code is just a base provided for your convenience.

- `bt_setup.[c|h]` : This is a C library written by the original author of this assignment that provides functionality for parsing command line arguments and peer strings.
- `bt_lib.[c|h]` : This is the core BitTorrent client library and header file. It has outlined a number of functions that you may need to write to complete this lab, with descriptions of their intended actions. Additionally, this file contains relevant structures for messages and arguments to your BitTorrent client.
- `bt_client.c` : This is the main portion of the client.
- `sample_torrent.torrent` : A sample torrent file to use in your development.
- `download_file.file` : The file that will be downloaded and uploaded by your client.

## 3 Non-BitTorrent Functionality

In addition to the core BitTorrent functionality, there are other standard software engineering challenges that you will encounter. Below, I outline some of the challenges and provide some advice for addressing them.

**Program Command Line Arguments:** I have built in basic command line parsing procedures into your client. Below is the usage:

```
bt-client [OPTIONS] file.torrent
-h          Print this help screen
-b ip       Bind to this ip for incoming connections, ports
            are selected automatically
-s save_file Save the torrent in directory save_dir (dflt: .)
-l log_file Save logs to log_file (dflt: bt-client.log)
-p ip:port  Instead of contacting the tracker for a peer list,
            use this peer instead, ip:port (ip or hostname)
            (include multiple -p for more than 1 peer)
-I id       Set the node identifier to id, a hex string
            (dflt: computed based on ip and port)
-v          verbose, print additional verbose info
```

**Program Output:** You should manage your program output such that information printed to the screen is useful but not overwhelming. Recall that your client will have a verbose flag, and you should reserve overly detailed output for verbose settings; however, there is a standard set of output you must provide. ***Your client must, at minimum, output the download status of the torrent, the number of connected peers, the amount of data uploaded, and the amount of data downloaded.*** Below is some sample output, note that you don't have to delete previous output.

```
File: foobar.mp3 Progress: 12.8% Peers: 5 Downloaded: 12 KB Uploaded: 15 KB
File: foobar.mp3 Progress: 13.1% Peers: 5 Downloaded: 14 KB Uploaded: 15 KB
```

Additionally, you should not print continually to the terminal. This will greatly slow down your program. Instead, you should print at reasonable intervals, for example, once a piece completes.

**Multiple Socket Connections:** Your client must be able to handle multiple connections seamlessly. That is, a client will have open sockets, and it needs to be able to read and send on all of them. Note that some operations are blocking, such as reading from a socket when there is no data to be read. Refer to the manuals for more information about how to use non-blocking procedures with sockets, but **I strongly recommend that you use the `select()` and `poll()` which both check the status of a file descriptor (like a socket) to see if it is ready for reading/writing/etc.**

**File I/O:** The core feature of BitTorrent is the ability to exchange sub-parts of the file in any order. This means you need to be careful about managing your file input and output to ensure that you are writing to the right part of the file based on the piece currently being downloaded. I refer you to the `netcat_part` project for detailed information about file seeking and file streams.

**Hashing:** BitTorrent makes use of SHA1 hashes to check data integrity. The hash, or digest, is a 20 byte string, and is easily computed using the OpenSSL library implementation of SHA1. In Figure 2, I provide some sample code for computing the SHA1 digest of data read from a file.

**Logging Information:** ***Your development must provide reasonable logging output.*** Events that must be logged are peer handshakes, peer messages, and file download status. You may also log other events, but describe these logs in your README. All log messages should contain: a *timestamp* in seconds/milliseconds since the start of the client; an *identifier* (in all caps) indicating the type of the log message, e.g., a peer



```

char data[BUF_SIZE]; //some size
char digest[SHA_DIGEST_LENGTH]; //20 bytes
int fd = open(some_file); //some file
int bytes_read, i;

bytes_read = read(fd, data, BUF_SIZE); //read data
SHA1(data, bytes_read, digest); //compute hash

//print the hash
for(i=0; i<SHA_DIGEST_LENGTH; i++) {
    printf("%02x", digest[i]);
}
printf("\n");

```

Figure 2: Sample code for computing the hash of data read from a file.

handshake; and, *parameters* indicating other aspects of the logged action, e.g., which peer the handshake occurred with. For example:

```

[102.12] HANDSHAKE INIT peer:192.168.1.2 port:22 id:0xdeadbeaf
[103.34] HANDSHAKE SUCCESS peer:192.168.1.2 port:22 id:0xdeadbeaf
[105.23] MESSAGE REQUEST FROM peer_id:0xdeadbeaf piece:10 offset:0 length:256
[107.25] MESSAGE PIECE TO peer_id: 0xdeadbeaf piece: 10 offset: 0 length:256
( . . . )

```

## 4 Development Advice

This is a large and complex assignment, and you will need to carefully plan your development with your lab partner. Below, I outline some development hints that should help with your planning and testing.

- **Mini-Goals:** Although there are many moving parts to this assignment, there are a lot of independent portions. Think small first, and then think big. Get one thing working, and then move on to the next. Here are some mini-goals you should look towards: (1) Simple parsing of bencoded files; (2) Handshake protocol; (3) Managing multiple clients; (4) Reading/Writing piece blocks from files; and etc. So long as you write down all the goals and tackle them in some sane order, you will find your project will go much better.
- **Skeleton Code is Not Cannon:** There is no requirement that you must use every part of the skeleton, use the same data structures, nor use the function outlines. This development is yours, so make it yours. Note that in many situations you will find it advantageous to define your own functions, alter or adapt provide function outlines, or complete scrap portions of the provided code base.
- **Divide and Conquer:** Although I believe that there is much to gain from pair programming, but if you and your partner spend all your time programming together you will not complete this assignment. Instead, once you divide the lab into distinct goals, you should focus on accomplishing these tasks individual and then later work together to merge developments. This mimics real world development.
- **Consider Error Conditions:** When testing your program, you should be considering error conditions from the get go. Ask yourself: What can go wrong? For example, how would your program respond if a client crashed? What happens if the reading/writing of a file is out of sync and pieces or corrupted? Can you recover or must the client quit?

- **Debug Early and Often:** In addition to using standard debugging tools, like `gdb` and `valgrind`, you should also build in your own debugging facilities via the log file and the verbose output flag.

## 5 Grading

**Milestone 1** For milestone 1, we expect you to be able to parse the `.torrent` files for the values of interest you need for this assignment. Specifically, you should be able to parse out values related to the `info` part of the `.torrent` file: `length` (length of the file to torrent), `name` (the actual name of the file to torrent), `piece length` (the size that each piece of the file is broken into), and `pieces` which includes the 20 bytes of a SHA1 hash per piece. As a hint: you may want to look at the `bt_info_t` structure in `bt_lib.h` for a spot to store these values.

In order for us to easily test this milestone, we ask that you print out to the terminal the results of parsing the file by listing the length, name, and piece length of the file (e.g. this should be part of your verbose output).

After you can read the `.torrent` file, you should put into place a roadmap of what else you must do to finish this project. Break your goals into bite-sized chunks: handshake protocol, implement sending a request message, etc. You may submit this roadmap as a `.txt` file.

Only one person in the group needs to submit a roadmap and base code. However, failure to submit Milestone 1 will result in a 10% deduction on the final grade for this project. Please note that submissions do not need to be perfect, we are simply looking to see that you have started this project and have a concrete plan of what needs to be done in the next two weeks.

**Grading Breakdown** Below is a general grading breakdown that will be used for this assignment. Note that other issues will be considered in the final grade, such as programming style, logic errors, memory leaks, etc..

- (B/B+) **1-Seeder and 1-Leecher:** Build a client that can implement the standard client protocol such that a single seeder can transfer a file to single leecher.
- (A-/A) **1-Seeder and N-Leechers:** Build a client that can support a single seeder and multiple leechers, but the leechers are not sharing with each other. E.g. Seeder 1 has the entire file, while Leecher 1, Leecher 2, and Leecher 3 are all trying to download the file from Seeder 1 at the same time.
- (A/A+) **N-Seeders and N-Leechers:** Build a client that can download a file from multiple seeders. These seeders should all support having N-leechers contact them, just like in the previous A-/A tier.