

Lab 1: Socket Programming: `netcat_part`

Assignment due: Sept. 17 at 11:00pm

Overview

The goal of this lab is to familiarize yourself with application level programming with sockets, specifically stream or TCP sockets, by implementing a client/server socket application, `netcat_part`. Additionally, this lab will introduce you to advanced file manipulation in C/C++ using file pointers and file streams. This lab is the first in a sequence that result in a small implementation of a BitTorrent client.

Readings

The following readings will help you understand the concepts of this project or be helpful to implement your code. Please remember that any readings assigned here will be material that you should have a good grasp on for assignments or tests!

- Kurose and Ross: Chapter 2.7
- Donahoo and Calvert (Sockets book): Chapter 2 (TCP sockets) – code from here and the examples we provide should provide a solid foundation for socket programming!

Deliverables

In addition to the PDF `write-up`, your submission should minimally include the following files for each part of the lab.

- `netcat_part.c`
- `Makefile`
- `README`

Your `README` file should contain a short header containing your name, username, and the assignment title. The `README` should additionally contain a short description of your code, tasks accomplished, and how to compile, execute, and interpret the output of your programs.

Even though you are working in groups on this assignment, each student must submit their own `write-up`. You may not see your partner's written answers. The `write-up` PDF file should contain a short header containing your name, username, and the assignment title. It should also contain general information about the contents of each file that you submit. That is, a person not familiar with this assignment should have a clear understanding of the purpose of each submitted file after reading your `write-up`. Finally, the

`write-up` should include a ‘Credits:’ section, which lists all the people (e.g., students in the class or not, friends, colleagues, etc.) and resources (e.g., webpages, manuals, books, code, etc.) that you consulted for this assignment.

Your program should always exit gracefully. In addition, the code should be clean and well-commented. Most importantly, the submitted code must compile and run on the CS `Burrow` machines; e.g. `silos.cs.indiana.edu`.

Cheating and Plagiarism

Although it is okay to discuss assignments (verbally; no written exchanges) with other students, all solutions handed in must be your own. Your code must be a group effort between the two group members only (you are expected to write the code together). The `write-up` must be written individually; you may **not** read your project partner’s (or any other student’s) `write-up`.

Copying answers from friends or the Internet is prohibited. We will use different software packages to check for plagiarism. Furthermore, you must cite all sources used for problem solutions (i.e. websites, books, discussions with classmates, etc.) in the `write-up`. For more information please see Indiana University Code of Student Rights, Responsibilities, and Conduct: <http://www.indiana.edu/~code/>.

Submission Instructions

Only one partner needs to submit a tarball with your files in folders named *netcat_part* via `oncourse`. Both partners must submit their own `write-up` PDF.

netcat_part **partial file transfers**

In this assignment you will familiarize yourself with basic socket programming and file I/O primitives by implementing a form of `netcat` for *secure* partial file transferring, so dubbed `netcat_part`. First, a description of standard `netcat` is provided followed by a description of the lab requirements for implementing `netcat_part`.

Basics of netcat

The `netcat` program (or `nc` on the CS machines) is a simple networking program that connects to a remote server and sends input from `stdin` over the network to the remote server. It can also function as a server by opening a socket and listening for incoming connection, writing all received data to `stdout`. This is why it is called `netcat`, or the “network version of `cat`”. Here is an example of some of its usages:

- **netcat as a client:**

The most common usage of `netcat` is as a client; that is, as a program that connects to a remote server. Here is a standard example:

```
>echo -e "GET /akaizer/test-nc.txt HTTP/1.1\r\nHost: homes.soic.indiana.edu \r\n\r\n" |  
nc homes.soic.indiana.edu 80
```

This command will connect to the CS web server on port 80 and issue a standard GET request for a test document, i.e `http://homes.soic.indiana.edu/akaizer/test-nc.txt`. The CS web server will respond with the text of the document, which will be printed to `stdout` by `netcat`. This example makes use of a shell pipe “|” which directs the standard output of one program as the standard input of another, i.e., `echo`’s standard output of “GET ...” is `netcat`’s standard input. We use the `-e` of `echo` in order to allow the use of the CRLF characters required by HTTP/1.1 to have a valid request.

- **netcat as a server:**

`netcat` also functions as a simple server that will listen for a connection and write all received data to standard out. Consider the command below:

```
#> nc -l 127.0.0.1 6767
```

Here, `netcat` is provided one command line arguments: `-l`, which indicates that `netcat` should listen for incoming connections. You can test your `netcat` server by trying to connect to it using another instance of `netcat` (a client instance). In another terminal on the same host, issue the following command:

```
#> echo "Hello World..." | nc 127.0.0.1 6767
```

In the terminal where the `netcat` server is running, you should see “Hello World...” printed to the terminal output. Once the connection is closed by the `netcat` client, the `netcat` server will also close its connection and exit.

This is the basic functionality of `netcat`, and the more you use it, the more you’ll see how useful it is. For a full description of `netcat`’s functionality, refer to the manual page.

Basics of netcat_part

You and your lab partner, will implement a form of netcat for securely transferring parts of files. After, downloading the skeleton code, you'll find `netcat_part.c` and a usage function, whose output is duplicated below:

```
netcat_part [OPTIONS] dest_ip [file]
    -h                Print this help screen
    -v                Verbose output
    -m "MSG"          Send the message specified on the command line.
                     Warning: if you specify this option, you do not specify a file.
    -p port           Set the port to connect on (dflt: 6767)
    -n bytes          Number of bytes to send, defaults whole file
    -o offset         Offset into file to start sending
    -l               Listen on port instead of connecting and write output to file
                     and dest_ip refers to which ip to bind to (dflt: localhost)
```

Your task is to implement all the functionality in the help reference above. Below, is detailed descriptions and example usages.

- netcat_part as client

Using netcat_part as a client is very similar to netcat, except that instead of reading input from stdin via echo, input is provided in the form of a file, file, or a message enclosed in quotations, MSG. For example, here is the same netcat_part code for sending "hello world..."

```
> ./netcat_part -m "Hello World..." 127.0.0.1
```

Likewise, you can send entire files using netcat_part:

```
> ./netcat_part remote.server.com moby_dick.txt
```

In addition to this standard usage, netcat_part has the additional feature of being able to send parts of a file, the _part of netcat_part. Consider a large file, such as the text of *Moby Dick*, and you only want to send the first 100 bytes of the file to a remote server rather than the whole file. You can use this netcat_part command:

```
#> netcat_part -n 100 remote.server.com moby_dick.txt
```

And if you want to send the first 100 bytes after an offset of 256 into the file, you can use this similar command:

```
#> netcat_part -o 256 -n 100 remote.server.com moby_dick.txt
```

Which will send 100 bytes of the file starting at byte 256. That is, it will send bytes 256 through 355 of the file.

- netcat_part as server

When functioning as a server, netcat_part acts nearly the same as netcat except that information sent over the line is written to the file rather than stdout. For example:

```
#> netcat_part -l localhost output_file
```

will open a socket for listening on the localhost and write data to the output file. The user should also be able to set a port for listening:

```
#> netcat_part -p 1024 -l localhost output_file
```

Your version of `netcat_part` functioning as a server does not need to worry about the offset or the number of bytes it receives when writing to the output file. It can just open the output file for writing, truncating the file if it already exists, and begin writing data to the start of the output file.

- security in `netcat_part`

The final goal of your `netcat_part` program is to achieve a level of security. In order to do this, you will utilize the `openssl/hmac.h` library in order to create a message authentication code (MAC) that will act as an integrity check for the data you send.

For the MAC to work, two parties share a secret key (a shared key) – we provide a simple sample shared key in this program. Using this shared key, the client can create a *hash* of the data that is being sent. For this project, the data you send from the `client` to the `server` will need to have the corresponding message digest included with it so that the server can calculate its own digest and compare that to the data it received to verify it is the same as the client sent it (i.e. no man-in-the-middle can change the data without you knowing about it). You can either compute this digest on a per packet basis or on a per file basis [e.g. send the digest at the end].

In order to compile this code in C, you will need to add the following to your `gcc`: `gcc -lssl netcat_part.c`

A brief example of using HMAC is shown here:

```
unsigned char *HMAC(const EVP_MD *type, const void *key, int keylen,
                   const unsigned char *data, int datalen,
                   const unsigned char *hash, unsigned int *hashlen);
```

- `type` The message digest to use (`EVP_sha1()` is a possible choice)
- `key` A buffer that contains the key that will be used
- `keylen` The number of bytes in the key buffer that should be used for the key
- `data` A buffer containing the data that an HMAC will be computed for.
- `datalen` The number of bytes in the buffer that should be used.
- `hash` A buffer that has the computed message digest. It should be at least `EVP_MAX_MD_SIZE` bytes in length.
- `hashlen` A pointer to an integer that will receive the number of bytes of the hash buffer that were filled.

You would pass in the message to be hashed as the `data` variable. The hash value will be stored in the `hash` variable that you use. So, when you send to the server you will send `data + hash + hashlen`.

Hint for if the HMAC is sent per packet: In order to avoid some potential pitfalls in TCP, you will want to reduce the amount of *data* you send per packet by enough bytes so that you can fit the hash and `hashlen` in that message. If you were previously sending 1024 bytes of data before you implemented security, you would want to now send `1024 - hashlen - EVP_MAX_MD_SIZE` bytes of data, while the rest of the space would hold the hash and hash length

Hints to implement

In order to implement this project, you will want to divide your focus into several different chunks and tackle each one before moving on. One such combination is: implement the client that sends messages, extend the client to send a file over sockets, implement the server, and finally implement the security component (HMAC + a function to compare two keys which can be compared by $k1[i] \neq k2[i]$). As long as you and your partner both agree to how to implement this assignment, things should go well!

File Programming Preliminaries and Socket Programming

Below are descriptions of the requisite C functions you will need to complete this assignment. First, file I/O is discussed, particularly file streams, and following, the basics of socket programming is discussed, including C-style pseudo-code examples. You can find I/O examples for C++ online (for example: <http://www.cplusplus.com/doc/tutorial/files/>), but the socket programming examples provided here work for both C/C++. More networking examples can be found online using Beej's guide <http://www.beej.us/guide/bgnet/output/html/singlepage/bgnet.html>.

File Streams

- **Opening a file:** You are probably already familiar with the basic file opening procedure `open()`, which returns a file descriptor, an `int`. Additionally, there are other ways for manipulating files in C using a file pointer. Consider `fopen()` below:

```
FILE * fopen(const char * filename, const char * restrict mode)
```

`fopen()` opens a file named `filename` with the appropriate mode (e.g., “r” for reading, “w” for writing, “r+” for reading and writing, and etc.). The important part to consider is that a file pointer is returned rather than file descriptor. A file pointer allows you to interpret files as streams of bits with a read head pointed to some part of the file.

For example, if you open a file for reading, the read head will be pointed to the beginning of the file, and if you open a file for append, the read head is pointed to the end of the file. Note that this is the way that Python handles file I/O, so you may already be familiar with this form even if you weren't aware.

- **Reading and Writing:** Reading and writing from a file pointer is very similar to that of a file descriptor:

```
size_t fread(void * ptr, size_t size, size_t nitems, FILE * stream);  
size_t fwrite(void * ptr, size_t size, size_t nitems, FILE * stream);
```

Like before data is read from or written to a buffer, `ptr`, but the amount is described as `nitems` each of `size` length. This is very useful when reading chunks of data, but you can always set `size` to 1 and `nitems` to the number of bytes you wish to read and write, i.e., read `nitems` each 1 byte in `size`.

- **File Seeking:** Seeking in a file is the primary reason to use a file pointer. This is a procedure that allows you to move the read head to arbitrary locations in the file. Here is the core function:

```
int fseek(FILE * stream, long offset, int whence);
```

which moves the read head of the file stream to the `offset` (measure in bytes) from a starting position whence. For example, to seek 100 bytes into a file, use this seek call:

```
int fseek(FILE * stream, 100, SEEK_SET);
```

Where `SEEK_SET` refers to the start of the file. You may find the following functions also useful: `ftell()`, which returns the current file position; and `rewind()`, which resets the file read position to the start of the file.

- **Converting to File Descriptor:** Finally, I should note that you can always convert a file pointer to a file descriptor and vice versa using either `fdopen()` and `fileno()`.

Socket Programming API

Note, that you will be using standard stream sockets, TCP connections, and not raw sockets, so you will not need sudo rights for this assignment.

- **Opening a socket for streaming:** To open a stream socket, you will use the following function call.

```
sockfd = socket(AF_INET, SOCK_STREAM, 0);
```

A call to `socket()` takes a socket domain, e.g., internet socket `AF_INET`, the type of the socket, `SOCK_STREAM` indicate a stream session or TCP, and a protocol, which is 0 for `SOCK_STREAM` over IP. `socket()` returns a socket file descriptor, which is just an integer.

- **Connecting a socket:** With stream sockets you must first connect with the destination before you can begin transmitting data. To connect a socket, use the `connect()` system call:

```
int connect(int socket, const struct sockaddr *address, socklen_t address_len)
```

which takes a socket file descriptor, a pointer to a socket address, and the length of that address. The return value of `connect()` indicates a successful or failed connection: You should refer to the manual for more details on error conditions.

Regarding socket addressing: you'll probably want to use the `sockaddr_in` form of `sockaddr`, which is the same size and has the following structure members:

```
struct sockaddr_in {
    u_char    sin_len;
    u_char    sin_family;
    u_short   sin_port;
    struct    in_addr sin_addr;
    char      sin_zero[8];
};
```

And an `in_addr` is a `uint32`, or a 32 bit number, like an `int`. You can convert string representation of internet addresses to `in_addr`'s using either `getaddrinfo()` or `gethostbyname()`.

Note that if your computer has multiple interface, you must first `bind()` your socket to one of the interfaces using the `bind()` system call. It has the following function definition:

```
int bind(int socket, struct sockaddr *address, socklen_t address_len)
```

where the address indicates which of the servers interfaces, identified by IP address, this socket should use for listening (or sending).

- **Listening and Accepting a Connection:** On the server end, a socket must be set such that it can accept incoming connections. This occurs in two parts, first it requires a call to `listen()`, and second, a call to `accept()`. The function definition for `listen()` is as follows:

```
int listen(int socket, int backlog)
```

Of course, `listen()` first argument is the socket that the will be listened on. The second argument, `backlog` indicates the number of *queued* or *backlogged* incoming connections that can be pending waiting on an `accept()` call before a connection refused message is sent to the connecting client.

The `accept()` function is the key server side mechanism of socket programming. Let's start by inspecting its function definition:

```
int accept(int socket, struct sockaddr * address, socklen_t * address_len)
```

Essentially, given a socket that is listening to incoming connections, `accept` **will block** until a client connects, filling in the address of the client in `address` and the length of the address in `address_len`.

The return value of `accept` is very important: It returns a *new socket file descriptor* for the newly accepted connection. The information about this socket is encoded in `address`, and all further communication with this client occurs over the new socket file descriptor. Don't forget to close the socket when done communicating with the client.

- **Reading and Writing:** Reading and writing from an stream socket occurs very much like reading and writing from any standard file descriptor. There are a number of functions to choose from, I suggest that you use the standard `read()` and `write()` system calls. Writing is rather simple:

```
ssize_t write(int filedes, void * buf, size_t nbyte)
```

which, given a file descriptor (the socket) and a buffer `buf`, `write()` will write `nbyte`'s to the destination described in the file descriptor and return the number of bytes written.

Reading from a socket is slightly more complicated because you cannot be certain how much data is going to be sent ahead of time. First consider the function definition of `read()`:

```
ssize_t read(int filedes, void * buf, size_t nbyte)
```

Similar to `write()`, it will read from the give file descriptor (the socket), and place up to `nbyte` into the buffer pointed to by `buf`, returning the number of bytes read. However, consider the case where the remote side of the socket has written more than the buffer size of bytes. In such cases, you must place the `read()` in a loop to clear the line. Here is some sample code that does that:

```
while(read(sockfd, buf, BUF_SIZE)){  
    //do something with data read so far  
}
```

That is, the loop will continue until the amount read is zero, which indicates that there is no more data to read. But be careful, subsequent calls to `read()` when there is no data on the line will block until there is something to read.

Finally, a note about the manual pages for `read()` and `write()`, the relevant manuals are in section 2 of the manual, which you access this way:


```
#> man 2 read  
#> man 2 write
```

- **Closing a Socket:** To close a socket, you simply use the standard file descriptor `close()` function:

```
int close(int fildes)
```

Putting it all together

Given all of these functions above, it might be hard for you to see how everything works together. Below, we've provided some psuedocode for a server and a client to help that intuition.

Client:

```
int sockfd;
struct sockaddr_in sockaddr;
char data[BUF_SIZE];

//open the socket
sockfd = socket(AF_INET, SOCK_STREAM, 0);

//set socket address and connect
if( connect(sockfd, &sockaddr, sizeof(struct sockaddr_in) < 0){
    perror("connect");
    exit(1);
}

//send data
write(sockfd, data, BUF_SIZE);

//close the socket
close(sockfd);
```

Server:

```
int serve_sock, client_sock, client_addr_size;
struct sockaddr_in serv_addr, client_addr;
char data[BUF_SIZE];

//open the socket
sockfd = socket(AF_INET, SOCK_STREAM, 0);

//optionally bind() the sock
bind(sockfd, serv_addr, sizeof(struct sockaddr_in));

//set listen to up to 5 queued connections
listen(sockfd, 5);

//could put the accept procedure in a loop to handle multiple clients

//accept a client connection
client_sock = accept(sockfd, &client_addr, &client_addr_len);

while(read(client_sock, data, BUF_SIZE)){
    // Do something with data
}

//close the connection
close(client_sock);
}
```

Credits

This lab is © Adam Aviv, which we have modified for P-538.